

In user space, following a program that calls `getpid()`:
`getpid()` is declared in `user.h`.

- `int getpid(void);`

`getpid()` is defined in `usys.S`.

- `SYSCALL(getpid)`
- This file contains a macro that expands to:
 - o `#define SYSCALL(getpid)`

<code>.globl getpid;</code>	<code>#Declares getpid as a global symbol.</code>
<code>getpid:</code>	<code>#Entry point of getpid.</code>
<code>movl \$SYS_getpid, %eax;</code>	<code>#Put system call number in eax register.</code>
<code>int \$T_SYSCALL;</code>	<code>#Trigger a software interrupt, enter the kernel.</code>
<code>ret</code>	<code>#Return result to the caller of getpid.</code>

`$SYS_getpid` is defined in `syscall.h`.

- `#define $SYS_getpid 11`

`$T_SYSCALL` is defined in `traps.h`.

Int `$T_SYSCALL` triggers a software interrupt and begins preparing to change from user space to kernel space.

- CPU saves the current state and calls the interrupt handler.
- The interrupt handler for `T_SYSCALL` is a `vector64`.
 - o Found in `vector.S`, which is generated by `vectors.p1`.
- `Vector64` jumps to `alltraps` function (`trapasm.S`).
- `alltraps` creates a trapframe and calls `trap(struct trapframe *tf)`. (`trap.c`).
- `struct trapframe` (`x86.h`) saves the user-space registers and `tf->eax` contains the system call number (`SYS_getpid`).

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitwise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: trapasm.S

#include "mmu.h"

# vectors.S sends all traps here.
.globl alltraps
alltraps:
# Build trap frame.
pushl %ds
pushl %es
pushl %fs
pushl %gs
pushal
# Set up data segments.
movw $(SEG_KDATA<<3), %ax
movw %ax, %ds
movw %ax, %es

# Call trap(tf), where tf=%esp
pushl %esp
call trap
addl $4, %esp

# Return falls through to trapret...
.globl trapret
trapret:
popal
popl %gs
popl %fs
popl %es
popl %ds

[ Read 32 lines ]
^G Get Help      ^O WriteOut      ^R Read File     ^Y Prev Page     ^K Cut Text      ^C Cur Pos
^X Exit          ^J Justify       ^W Where Is     ^V Next Page     ^U UnCut Text    ^T To Spell
```

`trap(struct trapframe *tf)` first checks if trap was invoked by a system call.

- `if(tf->trapno == T_SYSCALL)`
 - o In this case, it was so function enters condition.
- trapframe `tf` is saved to the current process control block (PCB).
 - o `myproc()->tf = tf;`
- Calls `syscall()`.
- Then returns to `alltraps`.
 - o Which restores all the user registers from the stack.
 - o Returns to user space with `iret`.

```

settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitwise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: trap.c

}

//PAGEBREAK: 41
void
trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL){
        if(myproc()->killed)
            exit();
        myproc()->tf = tf;
        syscall();
        if(myproc()->killed)
            exit();
        return;
    }

    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        if(cpuid() == 0){
            acquire(&tickslock);
            ticks++;
            wakeup(&ticks);
            release(&tickslock);
        }
        lapiceoi();
        break;
    case T_IRQ0 + IRQ_IDE:
        ideintr();
        lapiceoi();
        break;
    }
}

```

— Check if interrupt was caused by a syscall.
 — Is the process still alive?
 — Save the current trapframe and call syscall().

^G Get Help ^O WriteOut ^R Read File ^Y Prev Page ^K Cut Text ^C Cur Pos
 ^X Exit ^J Justify ^W Where Is ^V Next Page ^U UnCut Text ^T To Spell

syscall() reads the syscall number in `eax` (11 for `getpid`), and calls `sys_getpid` by:

- Getting a proc pointer to the current process which has the syscall number in the `eax` register.
 - o `struct proc *curproc = myproc();`
 - o `num = curproc->tf->eax;`
- calls `syscalls[SYS_getpid]` by invoking `syscalls[num]()`.
- `sys_getpid` is defined to be `syscalls[SYS_getpid]` (`syscall.c`).
- Return value is saved in `tf->eax`.
 - o `curproc->tf->eax = syscalls[num]()`;
- Control is returned to `trap`.
 - o Which then returns to `alltraps`, restores registers, and returns to userspace.

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitwise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: syscall.c

[SYS_close] sys_close,
[SYS_getprocs] sys_getprocs,
};

void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}

_____ Get a pointer a struct containing info
          about the current process.
_____ num = system call number stored in eax.
_____ Is this a valid system call number?
_____ Call syscalls[num]() which is defined
          to be syscalls[SYS_getpid]. Return
          value is stored in the eax register.

^G Get Help      ^O WriteOut      ^R Read File     ^Y Prev Page     ^K Cut Text      ^C Cur Pos
^X Exit          ^J Justify       ^W Where Is      ^V Next Page     ^U UnCut Text    ^T To Spell
```

sys_getpid (**sysproc.c**) returns **myproc()**->pid.

- **myproc()** defined in **proc.c**.
 - o returns a **struct proc** defined in **proc.h**.
 - o calls **pushcli()** to disable interrupts to make sure process isn't rescheduled while retrieving data. **popcli()** to return schedule interrupt functionality.
- **myproc()** creates and returns a **struct proc** (**proc.h**) that contains info about the current process running on the cpu.
 - o Size of memory, state, pid, parent, etc.
- **pid** member of the **proc struct** is returned to **syscall()**.
 - o Which is then stored in **tf->eax**.

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitvise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: sysproc.c

int pid;

if(argint(0, &pid) < 0)
    return -1;
return kill(pid);
}

int
sys_getpid(void)
{
    return myproc()->pid;
}

int
sys_sbrk(void)
{
    int addr;
    int n;

    if(argint(0, &n) < 0)
        return -1;
    addr = myproc()->sz;
    if(growproc(n) < 0)
        return -1;
    return addr;
}

int
sys_sleep(void)
{
    // Call myproc() (returns
    // a struct proc) and
    // return the pid
    // member.

    ^G Get Help    ^O WriteOut    ^R Read File    ^Y Prev Page    ^K Cut Text    ^C Cur Pos
    ^X Exit        ^J Justify     ^W Where Is     ^V Next Page    ^U UnCut Text  ^T To Spell
```

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitvise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: proc.c

    if (cpus[i].apicid == apicid)
        return &cpus[i];
    }
    panic("unknown apicid\n");
}

// Disable interrupts so that we are not rescheduled
// while reading proc from the cpu structure
struct proc*
myproc(void) {
    struct cpu *c;
    struct proc *p;
    pushcli();
    c = mycpu();
    p = c->proc;
    popcli();
    return p;
}

//PAGEBREAK: 32
// Look in the process table for an UNUSED proc.
// If found, change state to EMBRYO and initialize
// state required to run in the kernel.
// Otherwise return 0.
static struct proc*
allocproc(void)
{
    struct proc *p;
    char *sp;

    Returns struct proc*
    filled with info about
    the current process
    such as pid, parent
    pid, size, etc.

    ^G Get Help    ^O WriteOut    ^R Read File    ^Y Prev Page    ^K Cut Text    ^C Cur Pos
    ^X Exit        ^J Justify     ^W Where Is     ^V Next Page    ^U UnCut Text  ^T To Spell
```

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitvise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: syscall.c

[SYS_close] sys_close,
[SYS_getprocs] sys_getprocs,
};

void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}

Now the eax register
in the trapframe of
the current process
holds the pid of the
current process.

Return to trap.
```

^G Get Help ^O WriteOut ^R Read File ^Y Prev Page ^K Cut Text ^C Cur Pos
^X Exit ^J Justify ^W Where Is ^V Next Page ^U UnCut Text ^T To Spell

```
settings.tlp - wmiller@linprog.cs.fsu.edu:22 - Bitvise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: trap.c

}

//PAGEBREAK: 41
void
trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL){
        if(myproc()->killed)
            exit();
        myproc()->tf = tf;
        syscall();
        if(myproc()->killed)
            exit();
        return;
    }

    trap complete. Return
    to alltraps.

    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        if(cpuid() == 0){
            acquire(&tickslock);
            ticks++;
            wakeup(&ticks);
            release(&tickslock);
        }
        lapiceoi();
        break;
    case T_IRQ0 + IRQ_IDE:
        ideintr();
        lapiceoi();
        break;
}
```

^G Get Help ^O WriteOut ^R Read File ^Y Prev Page ^K Cut Text ^C Cur Pos
^X Exit ^J Justify ^W Where Is ^V Next Page ^U UnCut Text ^T To Spell

Transition back to user space by restoring the registers in **alltraps** after calling **trap**:

- Before jumping back to user space, we restore the user registers from the stack after returning from trap call in **alltraps** (**trapasm.S**).
- Returns to userspace after calling **iret**.



```
settings.tlp - wmill@linprog.cs.fsu.edu:22 - Bitwise xterm - linprog1:~/COP4610/xv6-public
GNU nano 2.3.1 File: trapasm.S

# Call trap(tf), where tf=%esp
pushl %esp
call trap
addl $4, %esp

# Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp # trapno and errcode
    iret
```

Diagram illustrating the flow of control:

- The **trapret** section restores all registers from the stack.
- The **iret** instruction returns to user space.

Terminal footer:

^G Get Help	^O WriteOut	^R Read File	^Y Prev Page	^K Cut Text	^C Cur Pos
^X Exit	^J Justify	^W Where Is	^V Next Page	^U UnCut Text	^T To Spell

Back in userspace:

- Process ID number has been returned to the user space for the calling program.